

AWANTIKA SRIVASTAVA

Data Scientist | LLM | MLOps

+91-8920482037 | sawantika81@gmail.com | [LINKEDIN](#) | [Github](#)

PROFILE SUMMARY

Data Scientist with **2+ years of hands-on professional experience** in designing, building, and deploying **data-intensive, scalable AI systems**. Strong expertise in **Python**, end-to-end **ML/LLM pipelines**, and production deployment of **Large Language Models (LLMs)**. Experienced in building **backend services, data platforms, and inference systems**, collaborating cross-functionally to deliver reliable AI products. Adept at translating complex requirements into **production-ready, cost-efficient, and scalable solutions**.

CORE TECHNICAL SKILLS

- **Programming Languages:** Python, C, C++, SQL
- **Machine Learning:** Supervised & Unsupervised Learning, Regression, Classification, Clustering, Random Forest, Decision Trees, SVM, KNN, K-Means, XGBoost.
- **Deep Learning & AI:** Neural Networks, CNN, RNN, LSTM, Transformers (BERT), Transfer Learning, Model Fine-tuning.
- **Computer Vision:** Image Classification & Preprocessing, Object detection (YOLO, SSD, MobileNet, ResNet), Video Analytics, OpenCV
- **NLP:** Text preprocessing, Tokenization, Transformers, Sentiment Analysis, Transformer & Diffusion-based models, Chunking, NLTK, SpaCy
- **LLMs:** LLM (GPT, Gemini, LoRA, QLoRA), RAG, Prompt Engineering,
- **MLOps & Deployment:** Docker, AWS (EC2, S3), MLflow, CI/CD for ML, Model Versioning, Monitoring.
- **Backend & Systems:** FastAPI, Flask, REST APIs, SQL
- **Tools & Collaboration:** Git, GitHub, Experiment Tracking, Documentation, Cross-functional Collaboration.

EXPERIENCE

Data Scientist | PPS International Pvt. Ltd.

January 2024-Present

- Designed and developed **AI-powered backend services** using **Python and FastAPI**, serving ML and LLM models via RESTful APIs.
- Built and maintained **end-to-end ML and LLM pipelines**, covering data ingestion, preprocessing, training, inference, and monitoring.
- Deployed and managed containerized AI services using **Docker on AWS**, focusing on scalability, reliability, and cost efficiency.
- Implemented **LLM-based production workflows**, including prompt engineering, retrieval pipelines, and inference orchestration.
- Designed data processing pipelines to support **large-scale data ingestion, feature engineering, and model evaluation workloads**.
- Collaborated cross-functionally with product and engineering teams to translate requirements into **sustainable AI solutions**.
- Applied MLOps best practices such as **model versioning, experiment tracking, CI/CD, and performance monitoring**.
- Contributed to architectural decisions and best practices to improve **system scalability, fault tolerance, and observability**.

PROJECTS

Railway Driver Assistance System (RDAS) | Enterprise ML Project

- Designed and deployed a **real-time ML systems** for unsafe behavior detection using **CNN-based object detection** models.
- Implemented **end-to-end ML pipelines** from data ingestion to production inference.
- **Deployed** optimized models using **TensorFlow Lite** for continuous, **low-latency inference**.
- Developed a **Flask-based** web interface to display detections and recorded video clips in **real time**.
- **Deployed** the complete solution on **edge devices**, enabling reliable on-device inference and **real-time event recordings**.

Amazon Stock Price prediction | Applied ML Project

- **Built batch-oriented ML workflows** for time-series forecasting using **LSTM** models.
- Implemented data **preprocessing**, feature engineering, and model **evaluation pipelines** on **large historical datasets**.
- Designed **sliding-window based sequence generation** and trained LSTM model 5-year stock price data.
- Evaluated model performance using appropriate regression **metrics** and **trend-based analysis** for short-term forecasting.

YouTube Comments Sentiment Analyzer | link - <https://youtube-ai-analyzer-ndzqo6r2mepjrtsdjmwxal.streamlit.app/>

- Deployed transformer-based **NLP models** as **production-ready** services with REST APIs.
- **Fine-tuned** and served a **DistilBERT-based sentiment** classification model for large-scale text inference.
- **Built** and deployed an interactive **streamlit web application** to perform real-time **sentiment analysis** on YouTube comments.
- **Processed** high-volume text **data** with sub-second inference **latency** for real-time sentiment analysis.

CERTIFICATION

- IBM Data Science & AI Certification
- AWS Generative AI with Large Language Models
- OpenCV Computer Vision Certification

EDUCATION

IMS Engineering College, Ghaziabad

September - 2020

Bachelor of Technology (Electrical and electronics engineering)